



Deploy Machine Learning Pipeline on Google Cloud Platform

Cloud Computing Project

Andrea Ciccotti - Andrea Bernini - Donato Francesco Pio Stanco

Introduction

The scenario

We imagined the following real scenario:

- a user enters a series of data about a football match
 - the match date, the name of the home team, and the away team.
- Once these parameters have been entered, they will be processed and a prediction of the outcome of the match will be provided in the output, i.e., victory of the home team, draw, or away team.



Tools

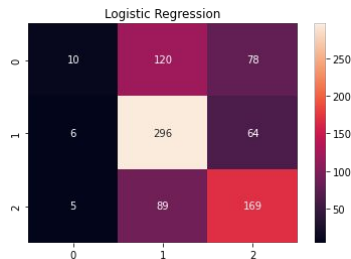
- **Flask**, for the creation of the web app through the use of Python, HTML, and CSS code.
- **Docker**, for the creation of a containerized environment in which to run our application.
- **Kubernetes**, for the management of containerized applications, or to completely manage their life cycle.
- **Google Cloud Platform (GCP)**, provides Google Kubernetes Engine which is an implementation of the open source Kubernetes framework, as well as various cloud computing services.



Implementation

Machine Learning Model

- We used a dataset that contains the last 10 season of Serie A.
 - 70 features and approximately 4000 rows.
- During the **preprocessing** phase:
 - Extrapolation of new features, using the existing ones, to avoid **overfitting**.
 - Transformation of the categorical features using a JSON to save the `string:number` combinations.
- We have tested several models locally (on Google Colab), and have chosen **Logistic Regression**.
 - it has the best trade-off between *Accuracy* and *Prediction Time*.

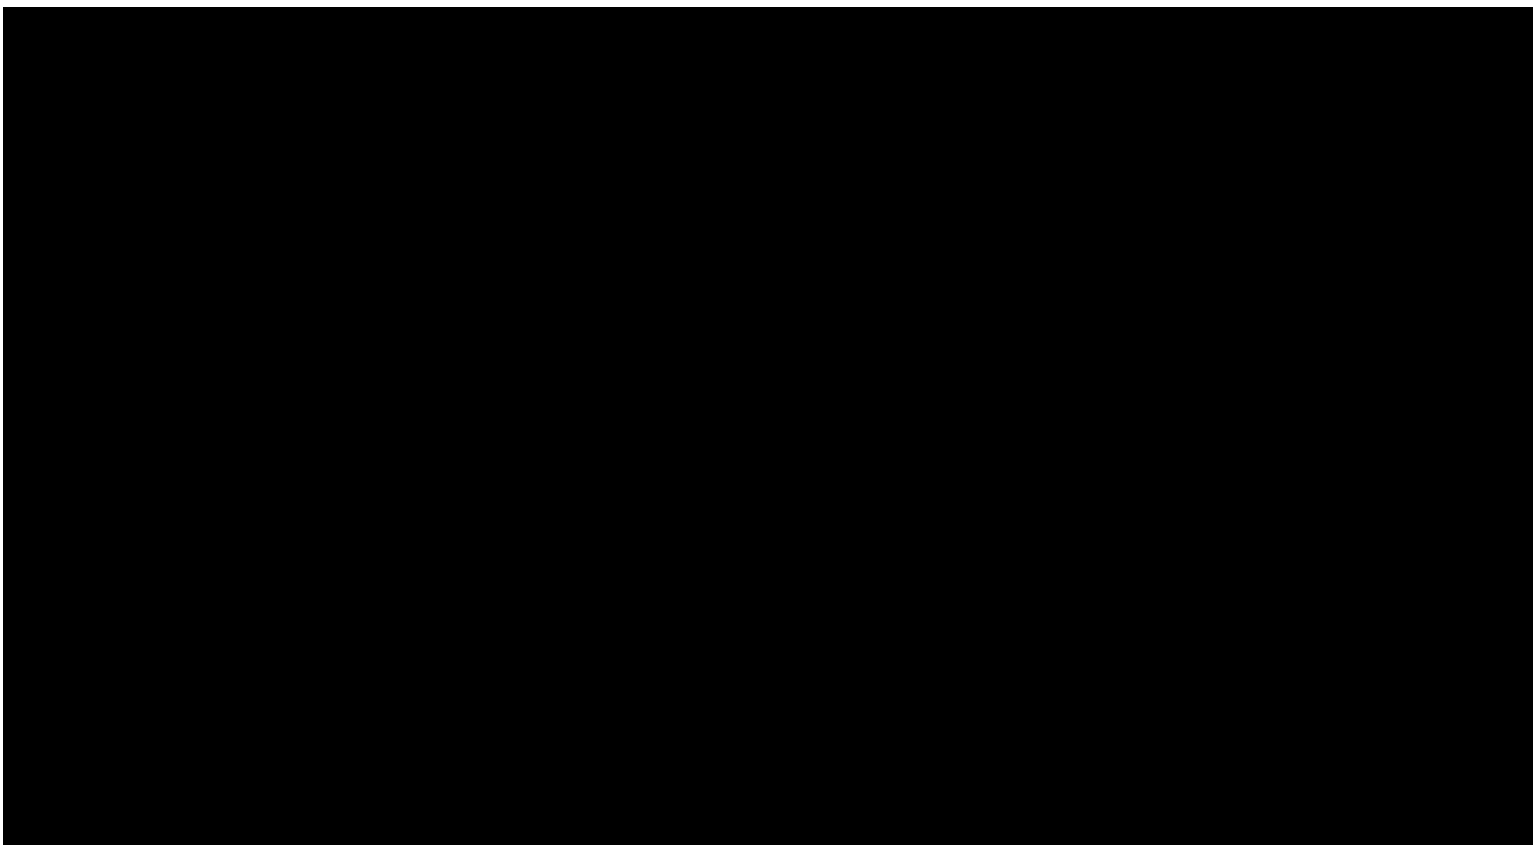


Web-App

Our web-app can be divided into two parts:

- **Front-end**, designed using HTML and CSS.
 - in which we have a FORM where the user can enter the data on which he wants to obtain a prediction.
- **Back-end**, developed using *Flask*, to render HTML code and manage HTTP Methods.
 - The main web-path of our application is `/predict` which allows to manage POST calls obtaining a prediction of the result.

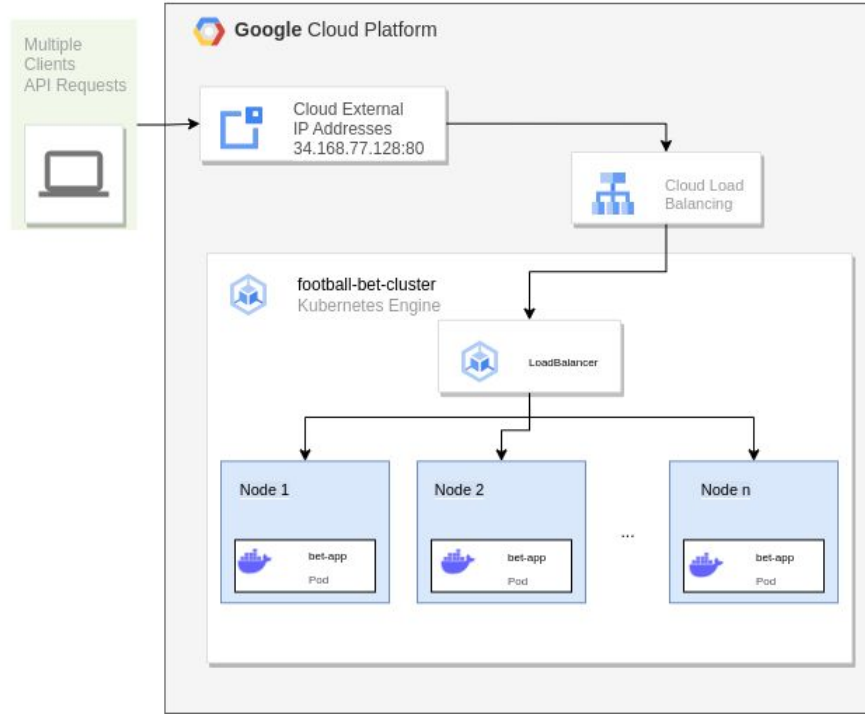




Deploy Web-App on GKE

1. Create a **Cluster** called *football-bet-cluster*, on **Google Kubernetes Engine**.
2. Build **Docker Image**, and upload it to **Google Container Registry**.
3. Deploy the image on *football-bet-cluster* in order to distribute and manage the application.
4. Finally expose the application, <http://34.154.93.188:80>.





Web App Architecture

Load Test

Simulation with Locust

We used Locust to perform load test:

- it gives you the ability to spread your performance tests across multiple machines.
- the architecture involves **two main components**:
 - Locust Docker container image
 - Container orchestration and management mechanism
- to perform load test we created a separate cluster called *loadtesting*



Start new load test

Number of users (peak concurrency)

Spawn rate (users started/second)

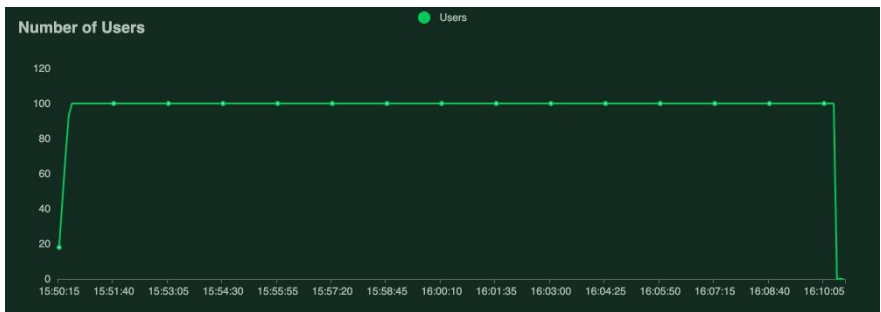
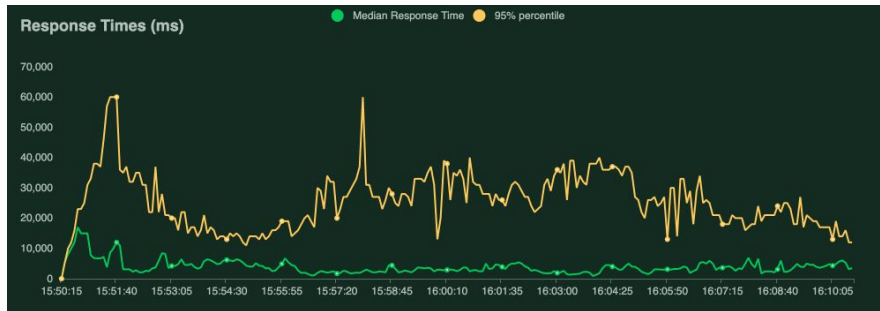
Host (e.g. http://www.example.com)

Start swarming

Test



- We have decided to apply **horizontal scaling**, adding more "machines" to our resource pool.
 - The goal is not to verify the correctness of the code or data, and furthermore, the responsiveness of the different page elements was not tested in the load/performance test.
- The path we tested for our application is `/predict_test` (managed back-end with Flask).
- As a case study we have analyzed a load of **100 users** and these are entered into the system with a spawn rate of **5 users/second**.
- The parameter we analyzed in the tests was the **CPU utilization**.
 - We did a series of tests to be able to understand what was the optimal number of pods to **obtain a lower response time** and also to **not have a large number of failures** during requests.



Request Statistics

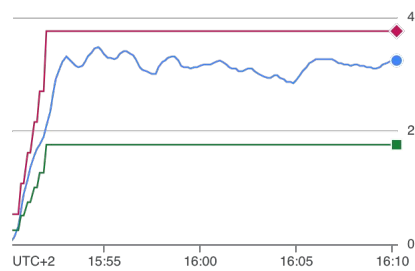
Method	Name	# Requests	# Fails	Average (ms)	Min (ms)	Max (ms)	Average size (bytes)	RPS	Failures/s
POST	/predict_test	12611	49	6473	1	60004	106	10.4	0.0
Aggregated		12611	49	6473	1	60004	106	10.4	0.0

Response Time Statistics

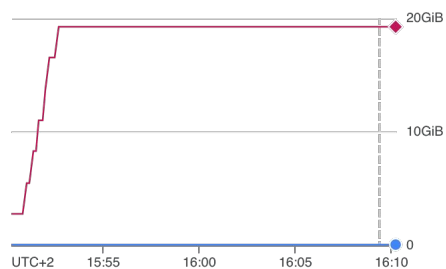
Method	Name	50%ile (ms)	60%ile (ms)	70%ile (ms)	80%ile (ms)	90%ile (ms)	95%ile (ms)	99%ile (ms)	100%ile (ms)
POST	/predict_test	3300	4900	6900	9700	15000	25000	44000	60000
Aggregated		3300	4900	6900	9700	15000	25000	44000	60000

Test with 7 pods: Locust

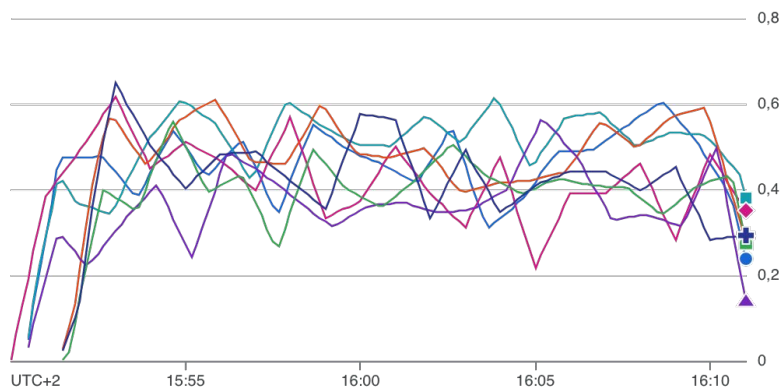
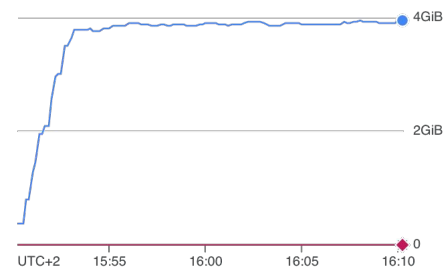
CPU



Disco



Memoria



Test with 7 pods: Google Resources

Estimated Bill





Costs of the project

- To analyze the costs of our project we used the **Google Cloud Price Calculator** service.
 - The cost of our project with the listed services is **\$225.31** per month for a total of **\$2703.72** for one year.
- Some aspects raises the price of our project:
 - the region chosen
 - the number of clusters used
 - the boot disk capacity of each node



Your Estimated Bill *

Estimated Monthly Cost: USD 255.31

CloudBuild	Cloud Build	1	USD 0.00
	Inbound data processed by load balancer	3.01361083984375e-06 GiB	USD 0.00
	Outbound data processed by load balancer	8.058547973632812e-06 GiB	USD 0.00
 4 x FootballBet cluster	e2-medium	2920 total hours per month	USD 61.64
 4 x boot disk	Persistent Disk - GKE Standard	100 GiB	USD 40.00
 3 x LoadTesting cluster	e2-medium	2190 total hours per month	USD 46.23
 3 x boot disk	Persistent Disk - GKE Standard	100 GiB	USD 30.00
GKE Clusters Fee	Zonal Clusters	1460 hours	USD 71.60
	Inbound data processed by load balancer	3.01361083984375e-06 GiB	USD 0.00
	Outbound data processed by load balancer	8.058547973632812e-06 GiB	USD 0.00
Used on standard VM instances IP addresses	Assigned and used in standard VM IPs	1460	USD 5.84
Total Estimated Monthly Cost			USD 255.31

Estimated Bill